

Image Recognition

Dr. Phil. Dativa Tibyampansha

Overview

- What is image recognition?
- What AI tools can be applied to achieve image recognition
 - Introduction to Convolutional Neural Networks (CNN)
- CNN Hands-on exercise

What is Image Recognition?

- Ability to see an image of an object and knowing what it is.
- We essentially classify everything that we see into categories based on their attributes
- This gives us ability to categorize items we have never seen before.
- Common AI applications

classification



Cat

object detection



Dog, Dog, Cat

segmentation



Dog, Dog, Cat

video processing, natural language processing

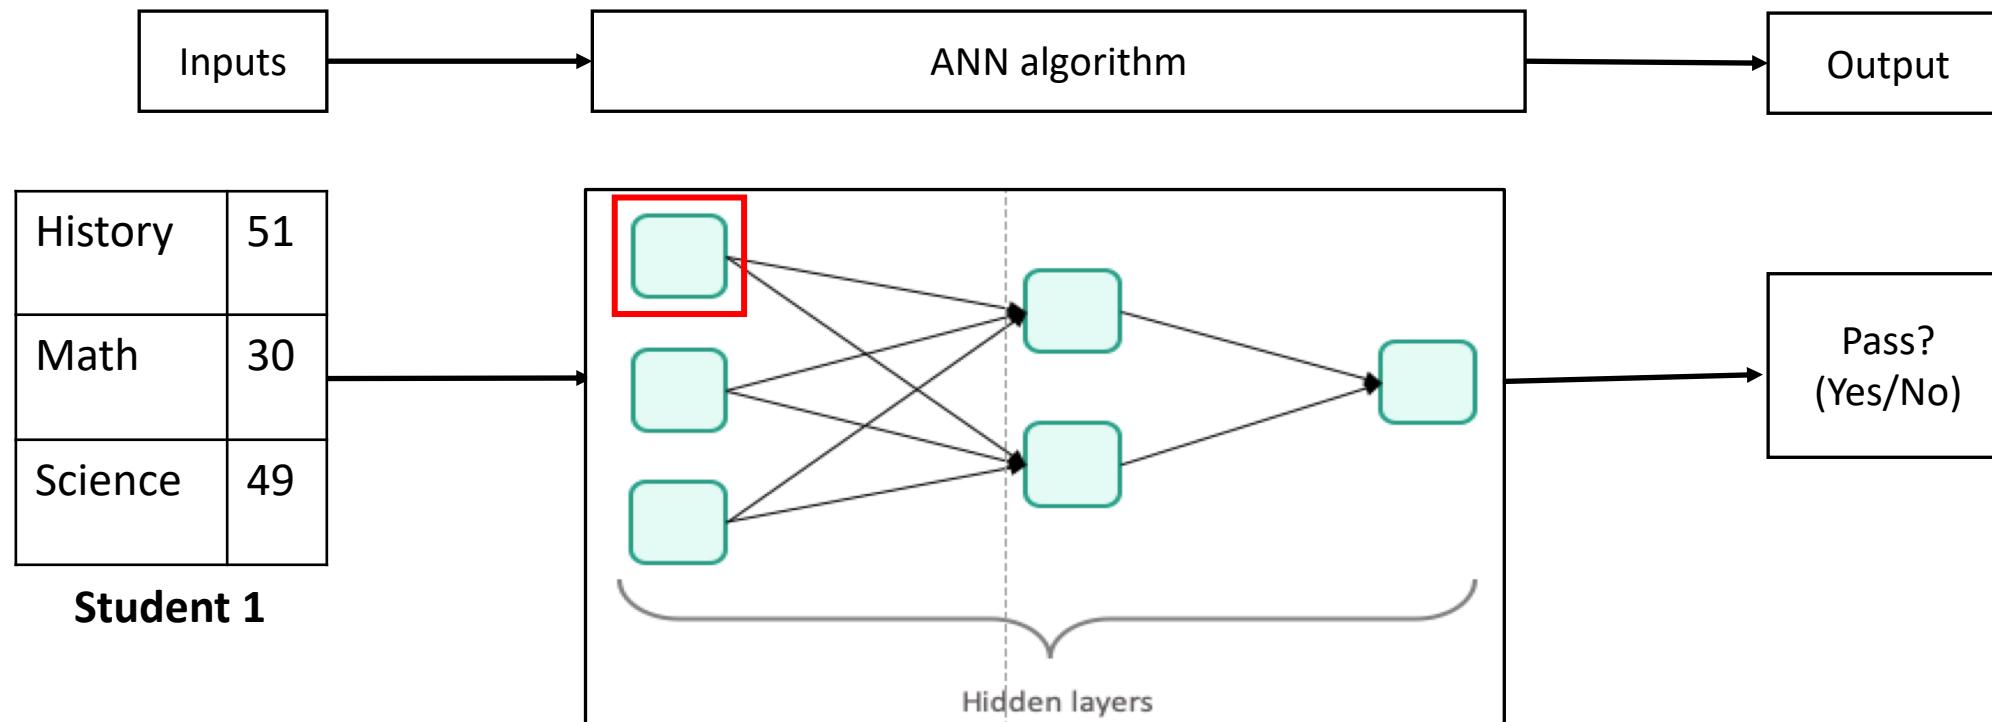
What AI for image recognition?

- Convolutional Neural Networks (CNN)
- Class of Artificial Neural Networks (ANN)

Artificial Neural Networks (ANN)- Intro

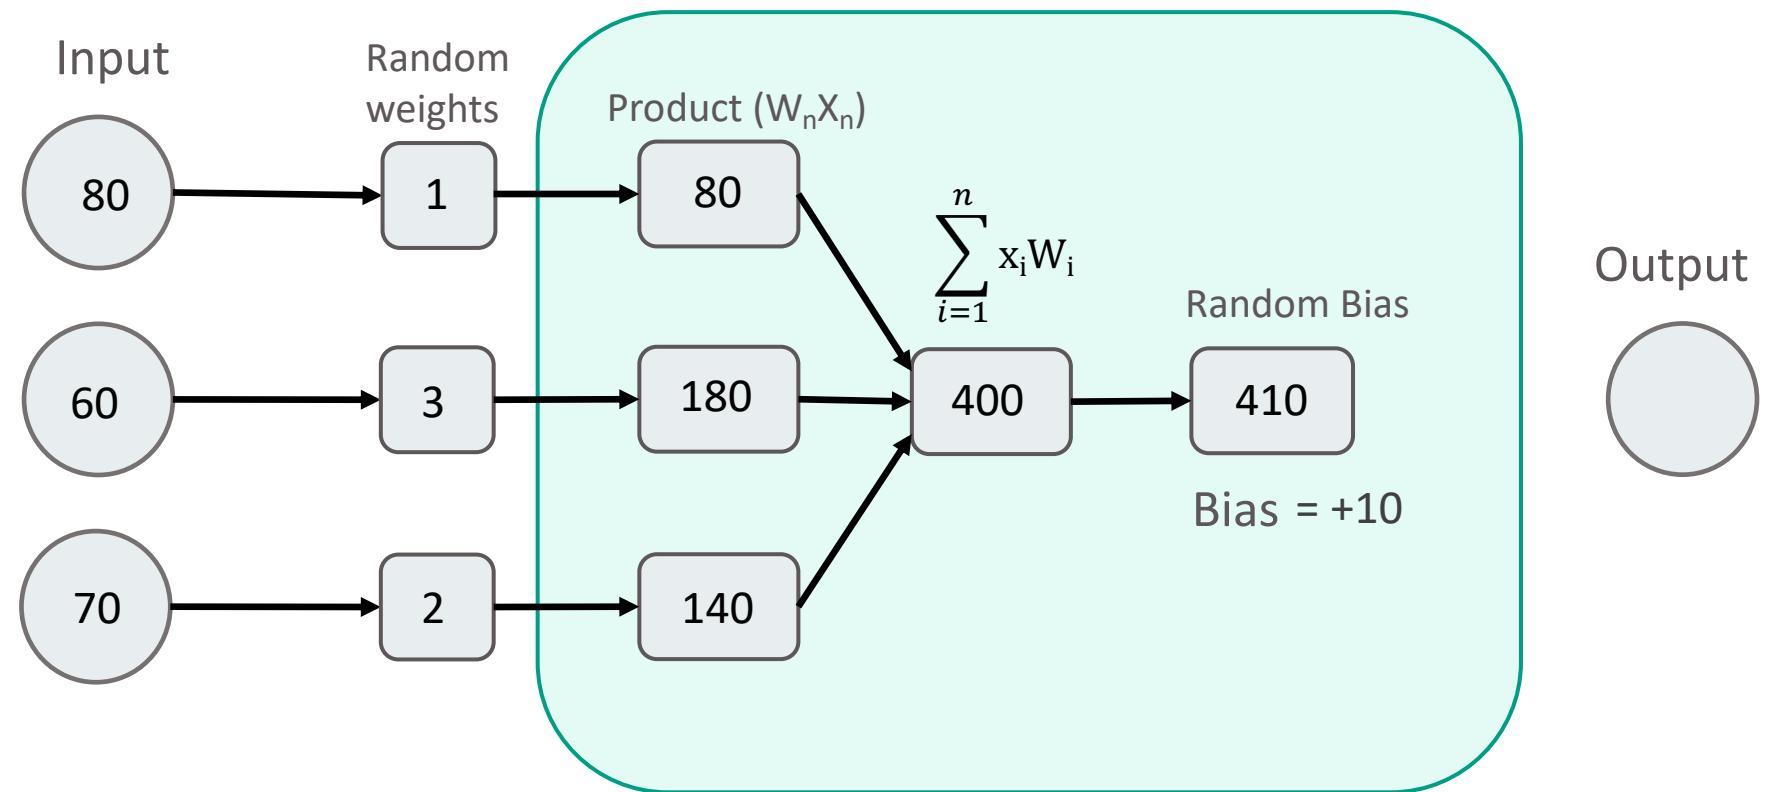
Case scenario: Predict if students will pass a final exam based on grades from previous students

Application: Early motivational guidance for students at risk of failing the final exam (“work harder”)



Neuron – info processing

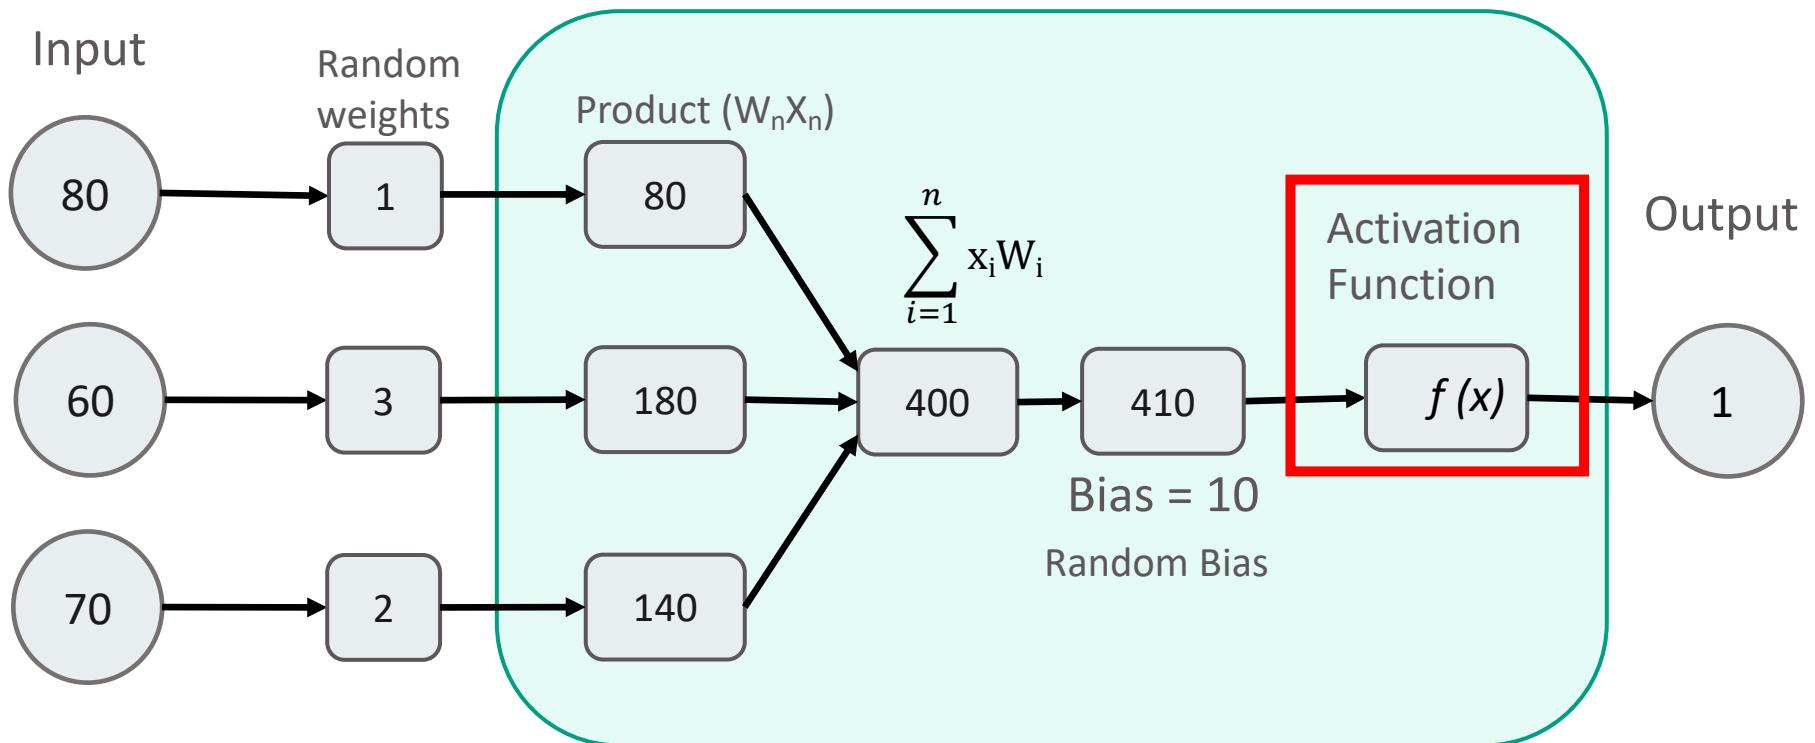
Student/ Subject	S3	S2	S1
History (H)	60	50	80
Math (M)	70	45	60
Science (S)	50	51	70



	S3	S2	S1
Grades	1 02.06.22	0	1

Neuron - Activation Function

Student/ Subject	S3	S2	S1
History (H)	60	50	80
Math (M)	70	45	60
Science (S)	50	51	70

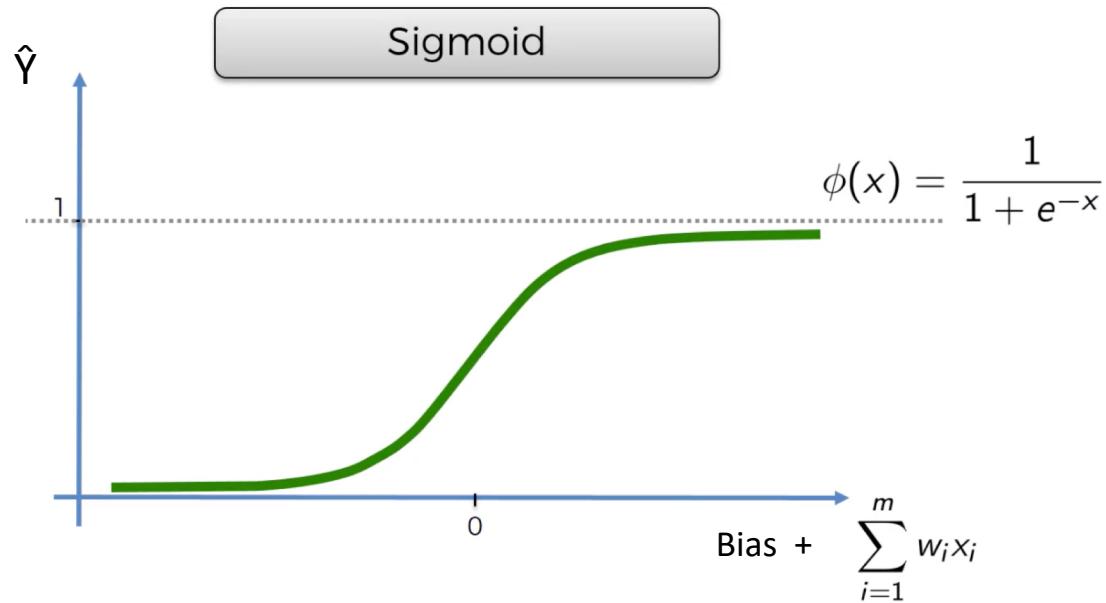


	S3	S2	S1
Grade	1	0	1

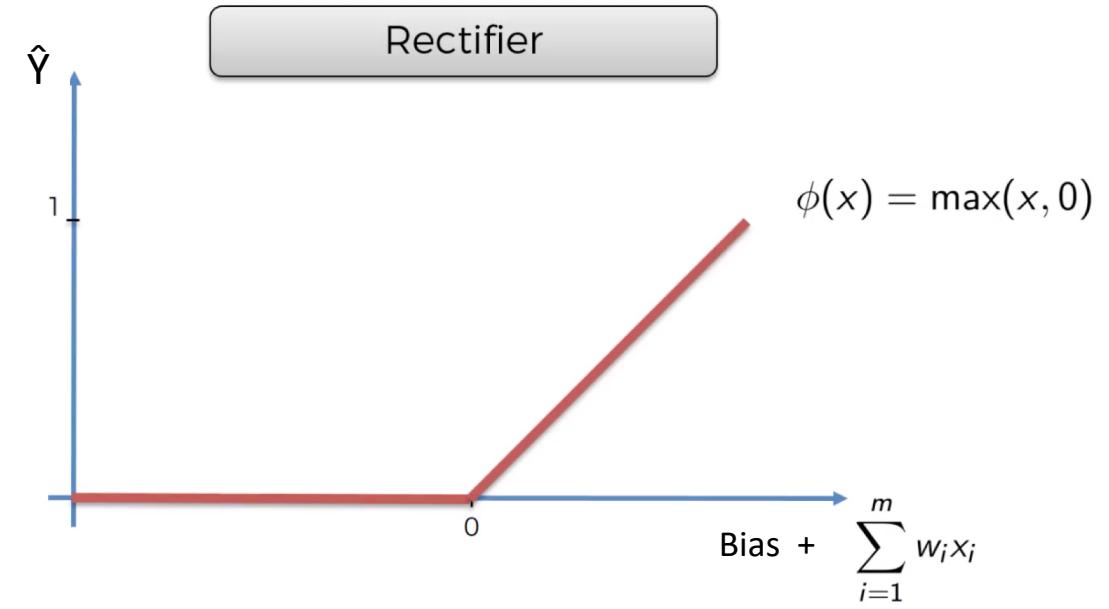
Activation Function: $f(x) = \begin{cases} 1 & \text{if } x \geq 200 \\ 0 & \text{if } x < 200 \end{cases}$

$$X = \text{bias} + \sum_{i=1}^n x_i W_i$$

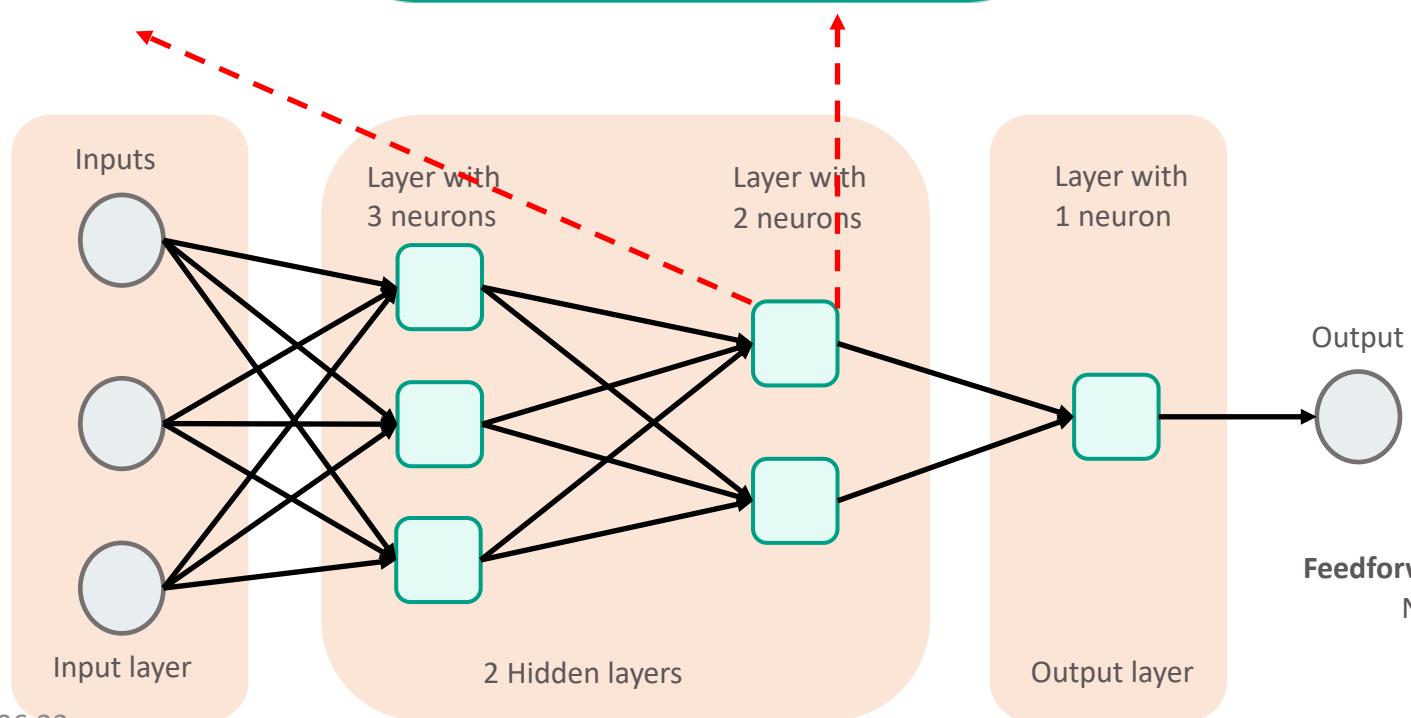
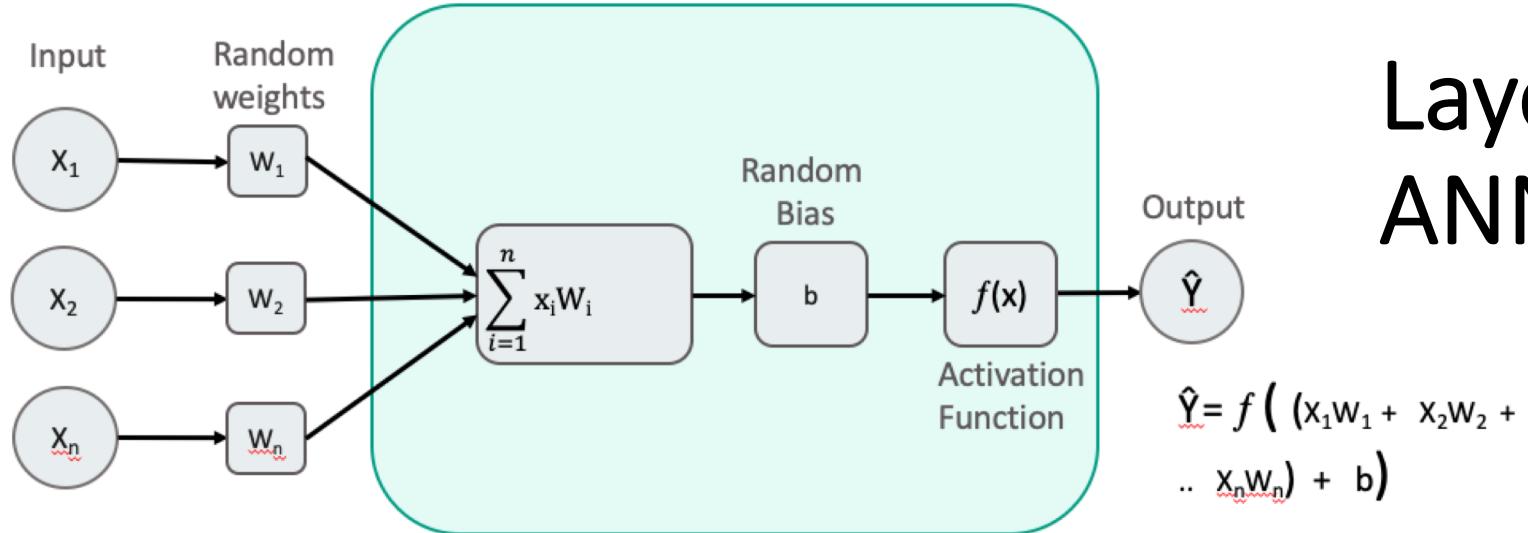
Examples of Activation Functions



\hat{Y} = predicted output

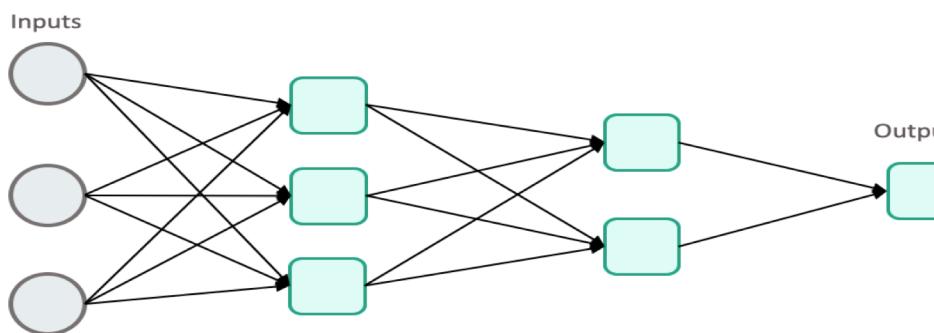


Layers – Higher-order ANN elements



How does an ANN learn?

	S3	S2	S1
H	60	50	80
M	70	45	60
S	50	51	70



	S1	S2	S3
Y	1	0	1

	S1	S2	S3
Ŷ	1	1	0

Cost function

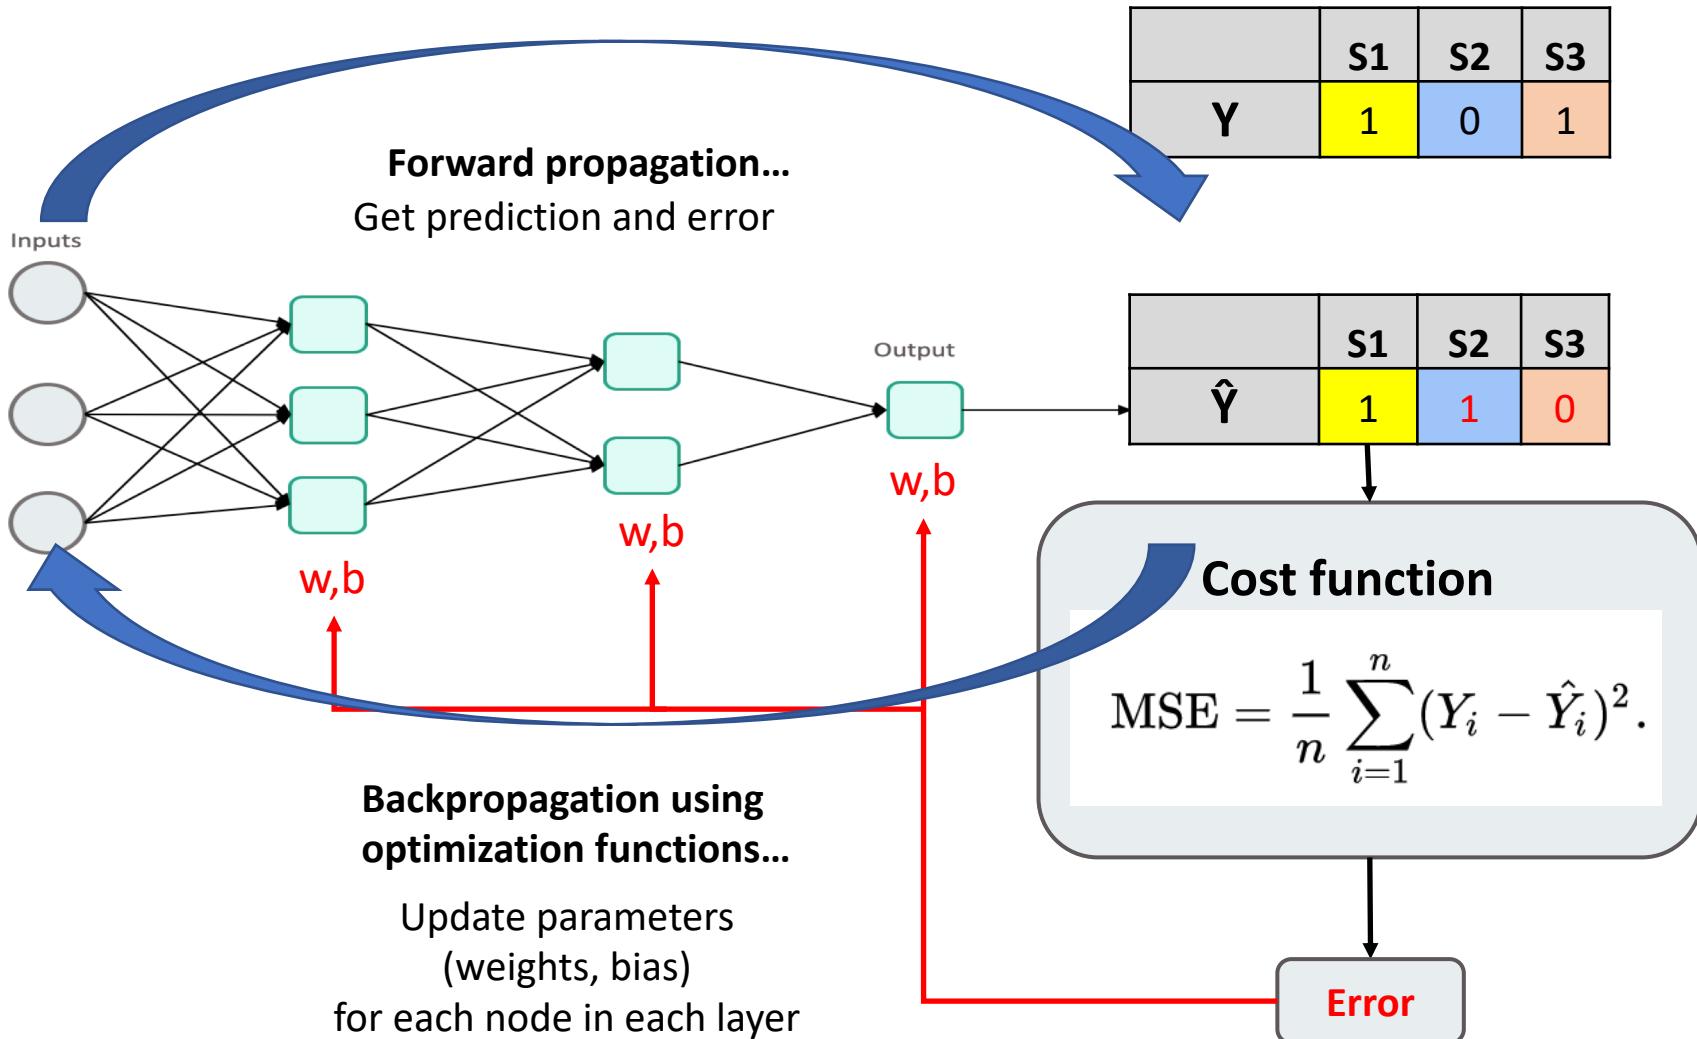
$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

- **Cost function:** measure of error between its estimated value (\hat{Y}) and true value (Y)
- Application and examples of cost functions
 - Regression: Mean Squared Error Loss,
 - Binary Classification: Binary Cross-Entropy, Hinge Loss
 - Multi-class Classification: Multi-class Cross Entropy Loss

<https://stats.stackexchange.com/questions/154879/a-list-of-cost-functions-used-in-neural-networks-alongside-applications>

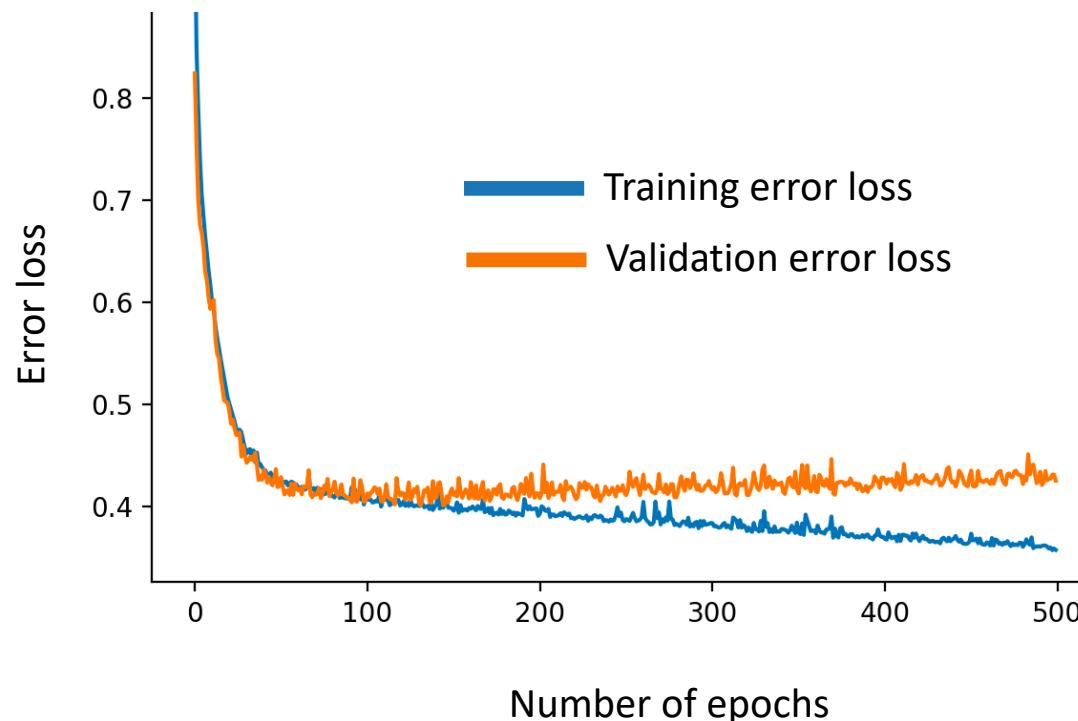
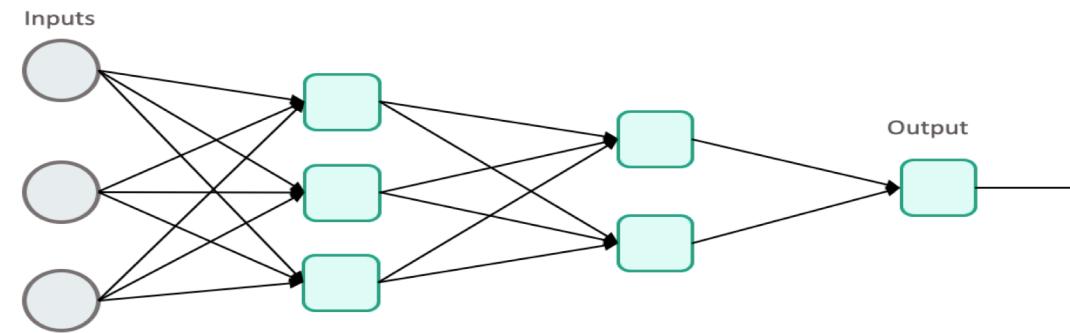
ANN – optimizing ...

	S3	S2	S1
H	60	50	80
M	70	45	60
S	50	51	70



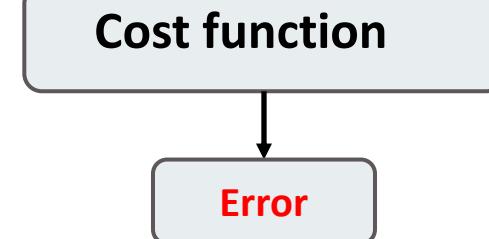
Validation

	S6	S5	S4
H	65	64	88
M	32	57	87
S	46	78	65



	S6	S5	S4
Y	1	0	1

	S4	S5	S6
\hat{Y}	1	1	1



Application of ANN

RNA Transcription Data (x)

	ENSG00000227232.4	ENSG00000233750.3	ENSG00000237683.5	
0	10.338030	0.362602	5.832498	
1	7.609902	0.361494	7.007413	
2	7.146308	0.194017	0.121261	
3	5.991603	0.207083	5.584339	
4	5.382497	0.256680	6.525889	
5	4.962599	0.069896	0.271817	
6	5.379276	0.273397	0.399012	
7	5.716319	0.103211	0.103211	
8	5.469149	0.106916	0.271401	
9	5.740473	0.364957	5.383119	
10	5.037652	0.456240	4.819037	
11	8.415640	0.117598	0.176398	
12	6.828088	0.228110	0.258524	
13	5.239712	0.243512	0.067176	
14	4.428489	0.119689	1.735489	
15	6.401946	0.235315	0.283762	
16	6.601959	0.153911	0.550838	
17	18.234970	0.227369	5.237059	

Labels(y)

Doxorubicin

0.106768

0.026413

0.160035

0.056119

0.035957

0.043444

0.017766

0.045722

0.026402

0.268859

0.020532

0.026486

0.045192

0.061365

0.010237

0.020460

0.034832

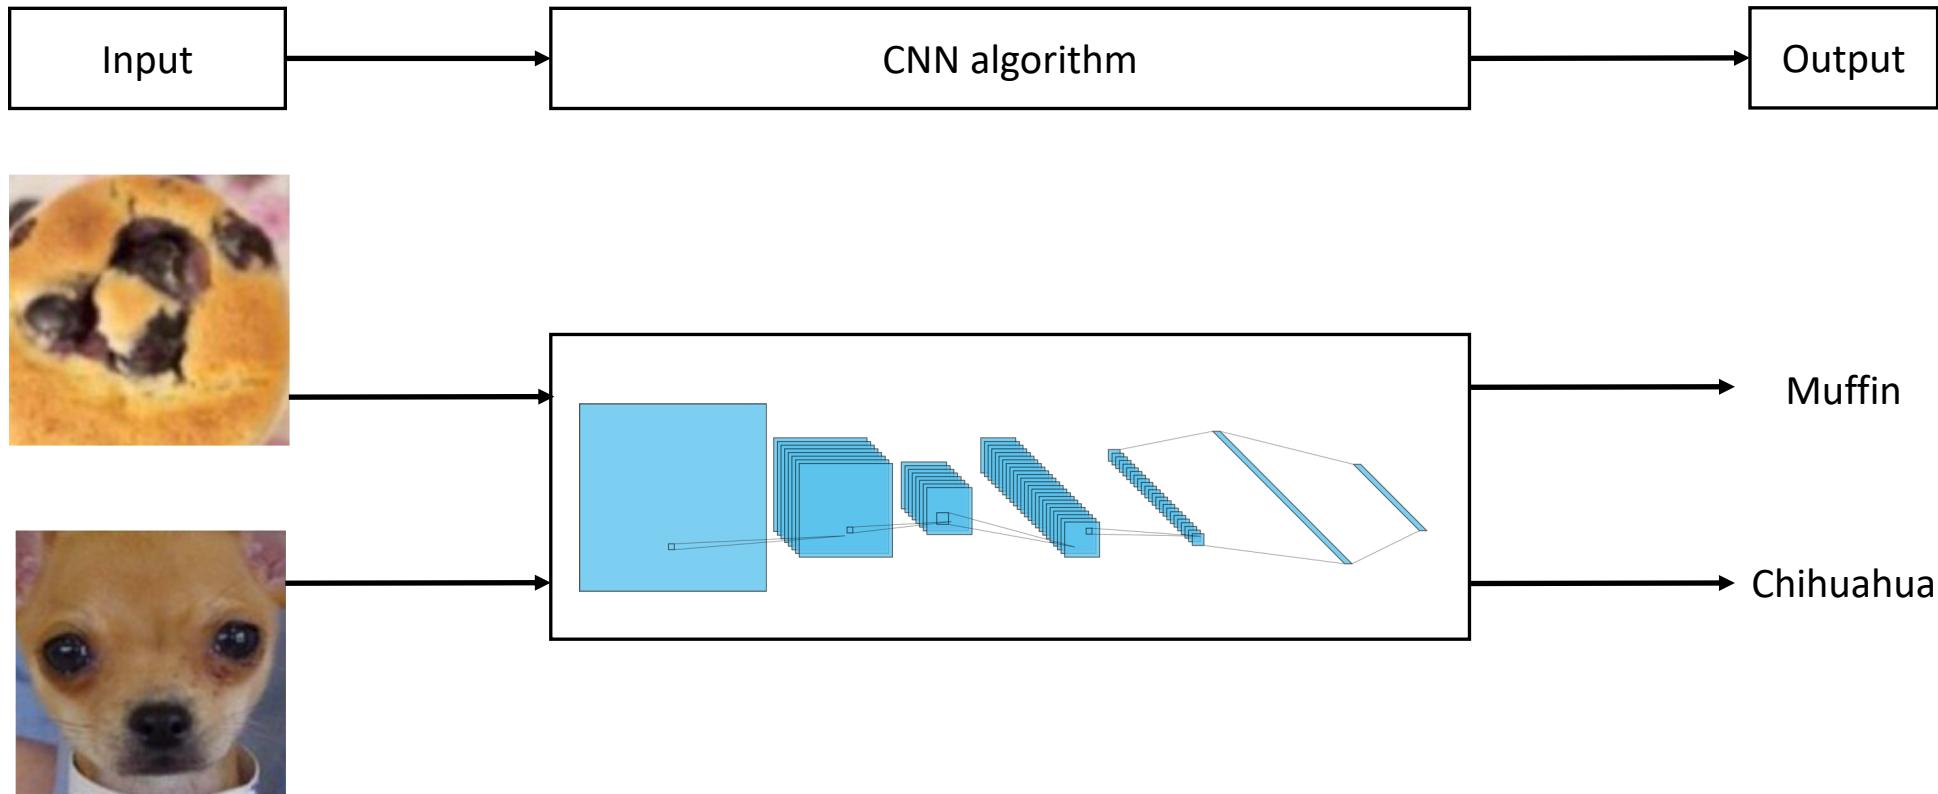
0.038035

ANN definition with Keras in python

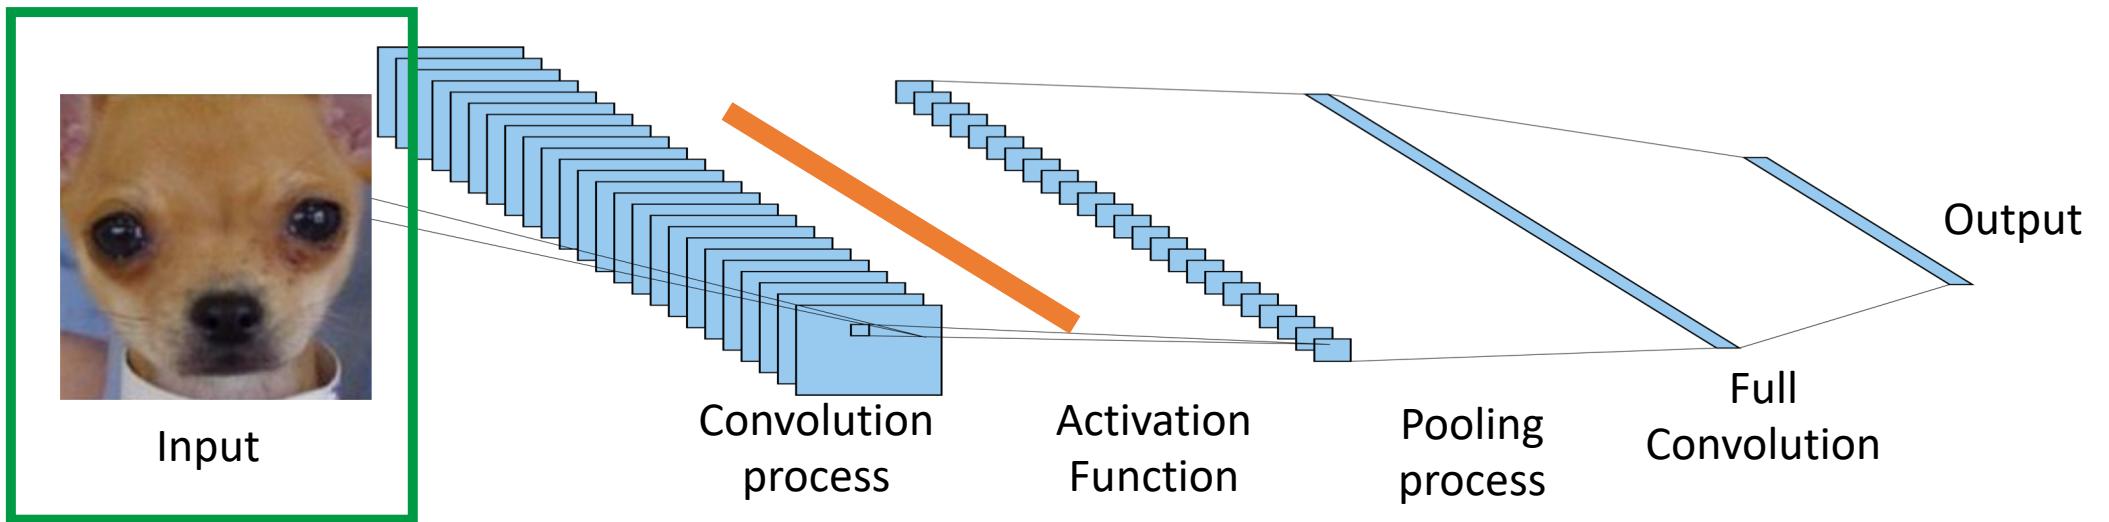
```
from numpy import loadtxt
from keras.models import Sequential
from keras.layers import Dense
# load the dataset
dataset = loadtxt('data.csv', delimiter=',')
# split into input (X) and output (y) variables
X = dataset[:,0:8]
y = dataset[:,8]
# define the keras model
model = Sequential()
model.add(Dense(12, input_dim=8, activation='relu'))
model.add(Dense(8, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
# compile the keras model
model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
# fit the keras model on the dataset
model.fit(X, y, epochs=150, batch_size=10)
# evaluate the keras model
_, accuracy = model.evaluate(X, y)
print('Accuracy: %.2f' % (accuracy*100))
```

Convolutional Neural Networks

Case scenario: Classify between a chihuahua and a muffin



CNN Steps



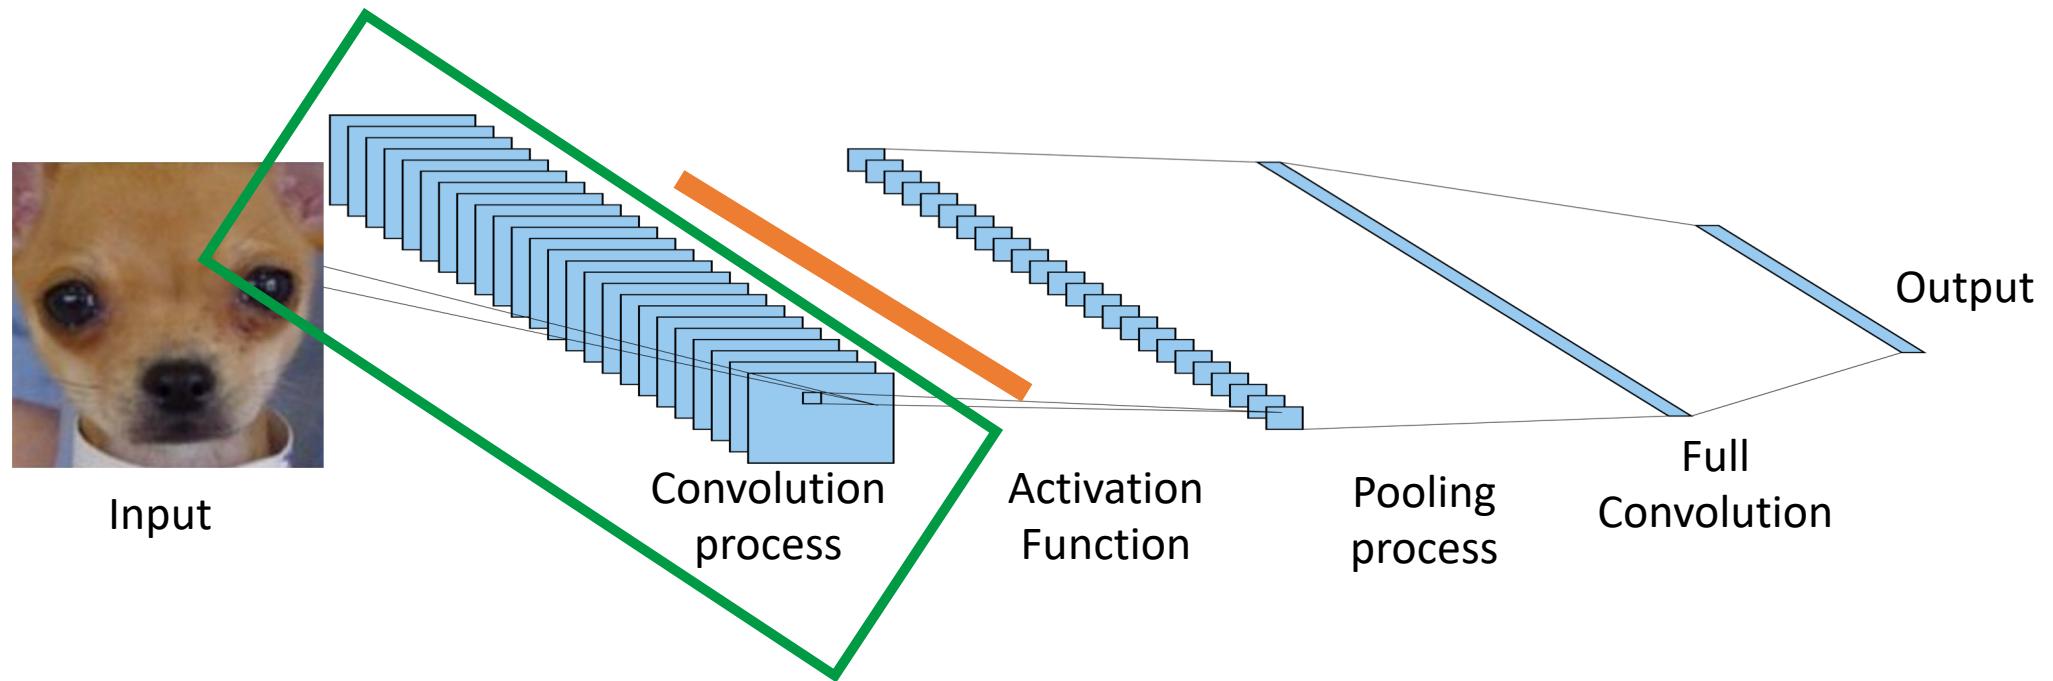
Input - images are numbers!

- Consist of pixels, with intensities 0 to 255
- Colors channels:
 - RGB (almost all) vs. BGR (OpenCV)
 - Other color spaces exist.
- Shape:
 - Computer vision: [height, width, channels] (3D)
 - Deep learning: [batch_size, height, width, channels] (4D)



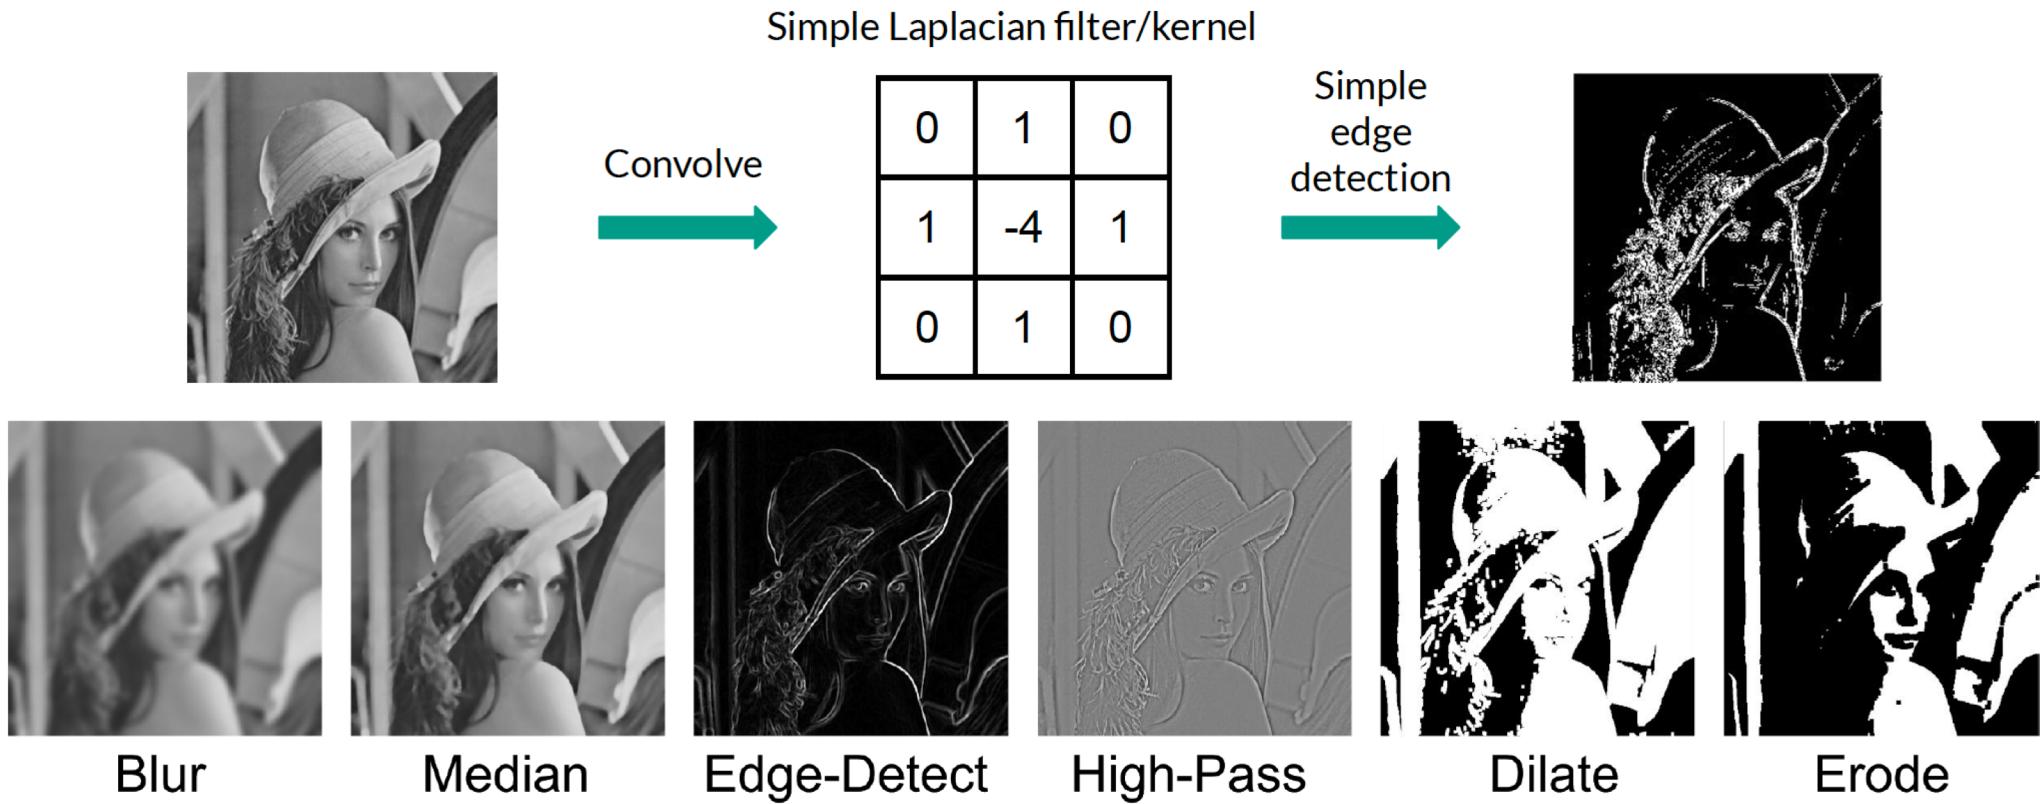
		Blue channel						
		Green channel			Red channel			
		171	200	19	6	...	26	
1	120	67	89	107	...	13	18	
2	12	216	145	26	...	181	81	
3	0	16	4	45	...	44	56	
4	0	78	90	167	...	25	...	
...	
64	12	67	82	141	...	12	12	
	1	2	3	4	...	64		

CNN Steps



Convolution

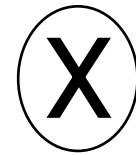
- matrix filters extract features from images



Convolution process

0	0	0	0	0	0	1
0	1	0	1	1	0	0
0	0	1	0	0	1	1
0	0	1	0	0	0	0
0	0	0	0	1	0	0
0	1	1	1	1	1	0
0	0	0	0	0	0	0

Image input

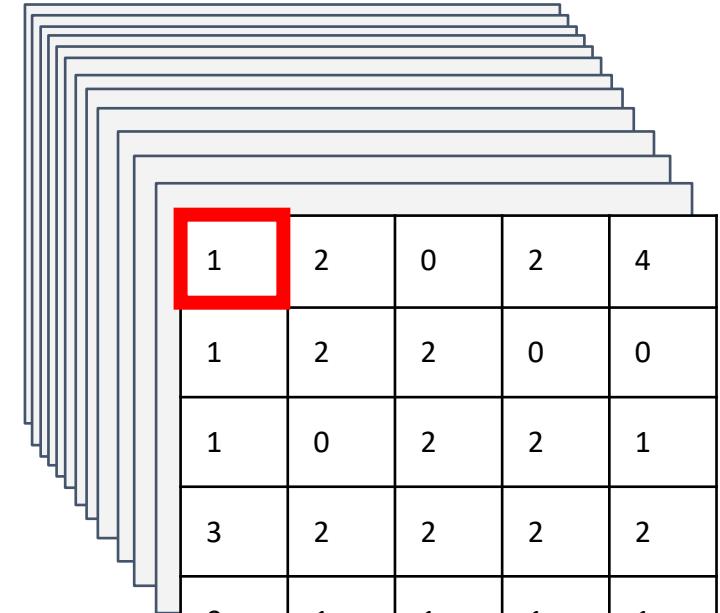


Random learnable
weights

0	0	1
1	0	0
0	1	1

Kernel (3X3)
Stride (1,1)

dot-product of image pixels
with a kernel weights



A stack of feature maps. The top-most feature map is shown in full, revealing a 5x5 grid of values. A red box highlights the value '1' at the top-left position (row 0, column 0). Below this, several other feature maps are shown as thin vertical slices, indicating the depth of the feature map volume.

1	2	0	2	4
1	2	2	0	0
1	0	2	2	1
3	2	2	2	2
0	1	1	1	1

Feature maps

Hyperparameters

- The size of the feature map can be controlled by the following hyper-parameters:

Stride: no. of pixels to skip when sliding over the input.

Diagram illustrating the convolution process. An image input (9x10 grid) is processed by a 3x3 kernel with stride 1. The output feature map (5x5 grid) shows the result of the convolution. A red box highlights the first step of the kernel's movement, and a red arrow points to the stride value (1,1).

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	0	0
0	0	0	1	0	1	1	0	0	0
0	0	0	1	0	0	1	1	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	1	0	0	0	0
0	0	1	1	1	1	1	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Image input

Padding: no. of zeros added around the input.

Diagram illustrating the convolution process. An input image (3x3 grid) is processed by a 3x3 kernel. The input has padding of 1 zero added around it. The output feature map (3x3 grid) shows the result of the convolution.

0	0	1
1	0	0
0	1	1

3x3 Kernel

Kernel size: sets the size of the filter

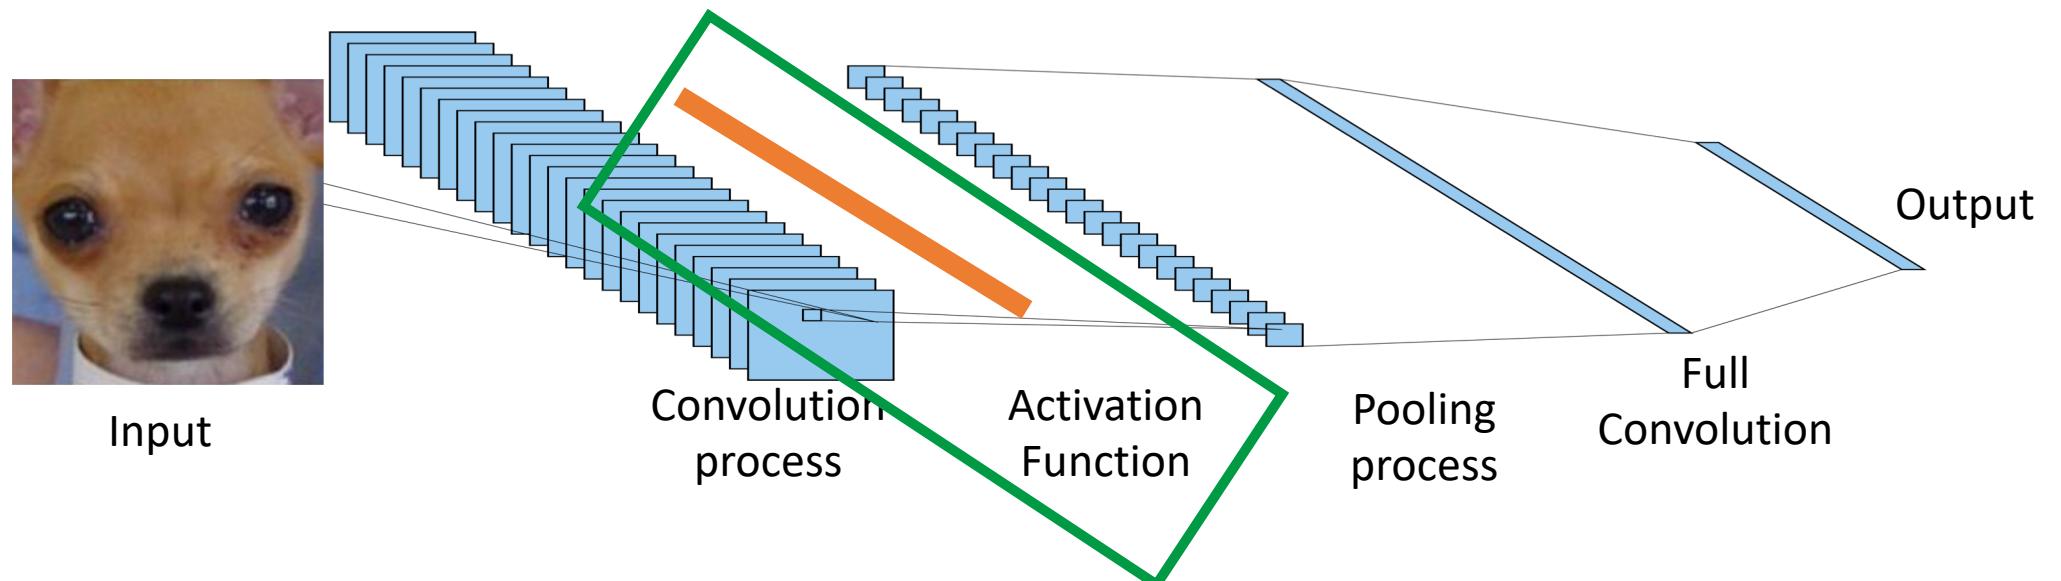
No. of filters: sets the depth of the output.

Diagram illustrating the convolution process. A single 3x3 kernel is applied to multiple input channels (represented by the stack of 3x3 grids) to produce a single feature map (5x5 grid). The resulting feature map has a depth of 5, corresponding to the number of filters.

1	2	0	2	4
1	2	2	0	0
1	0	2	2	1
3	2	2	2	2
0	1	1	1	1

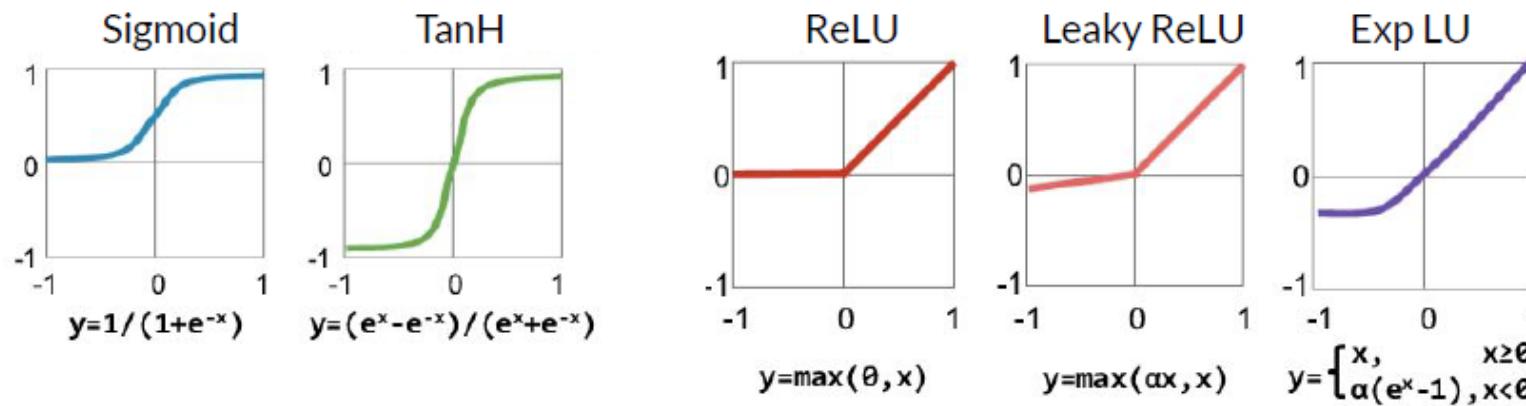
Feature maps

CNN Steps

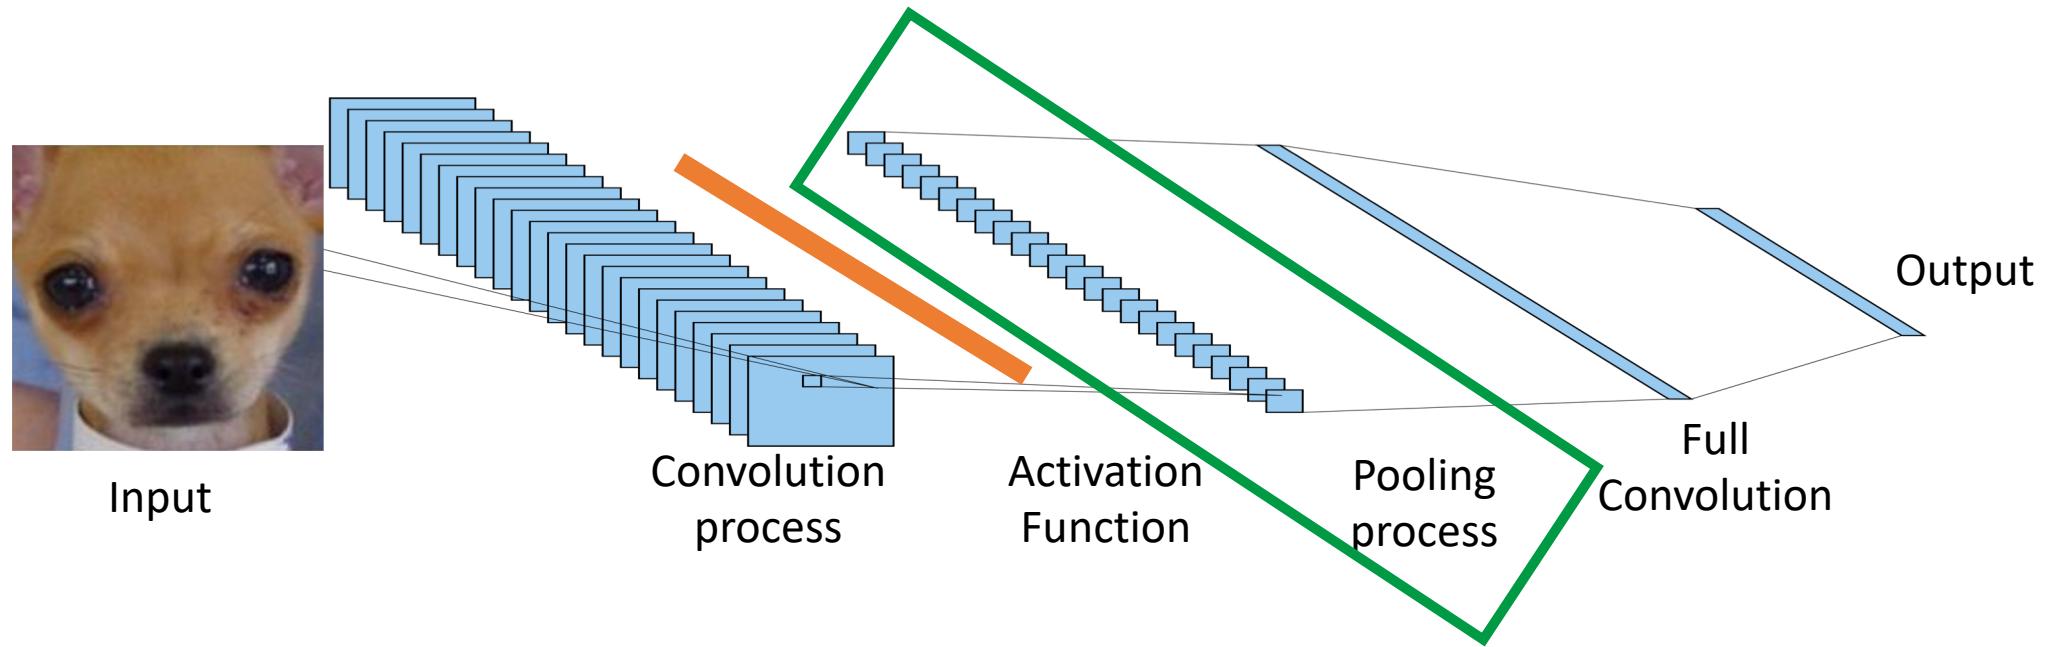


Activation Function

- Provide non-linearity to output.
- Applied to every feature map
- Rectified Linear units (**ReLU**) became popular in recent years.
 - They speed up training and help in convergence.
- Usually used directly after convolution.

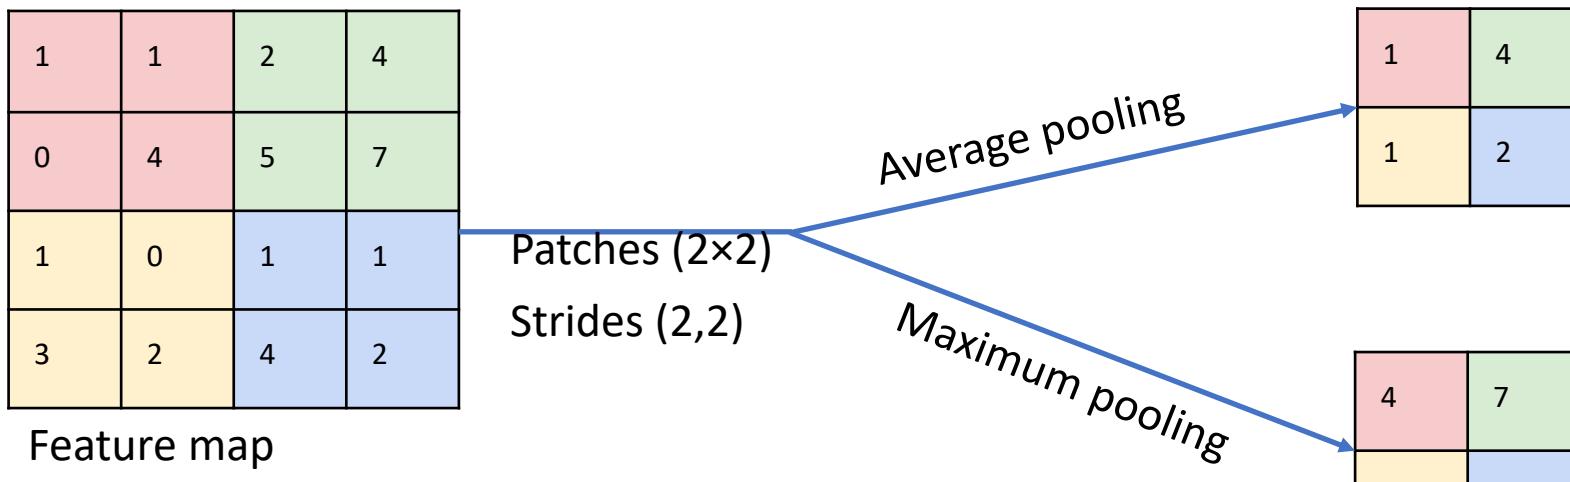


CNN Steps

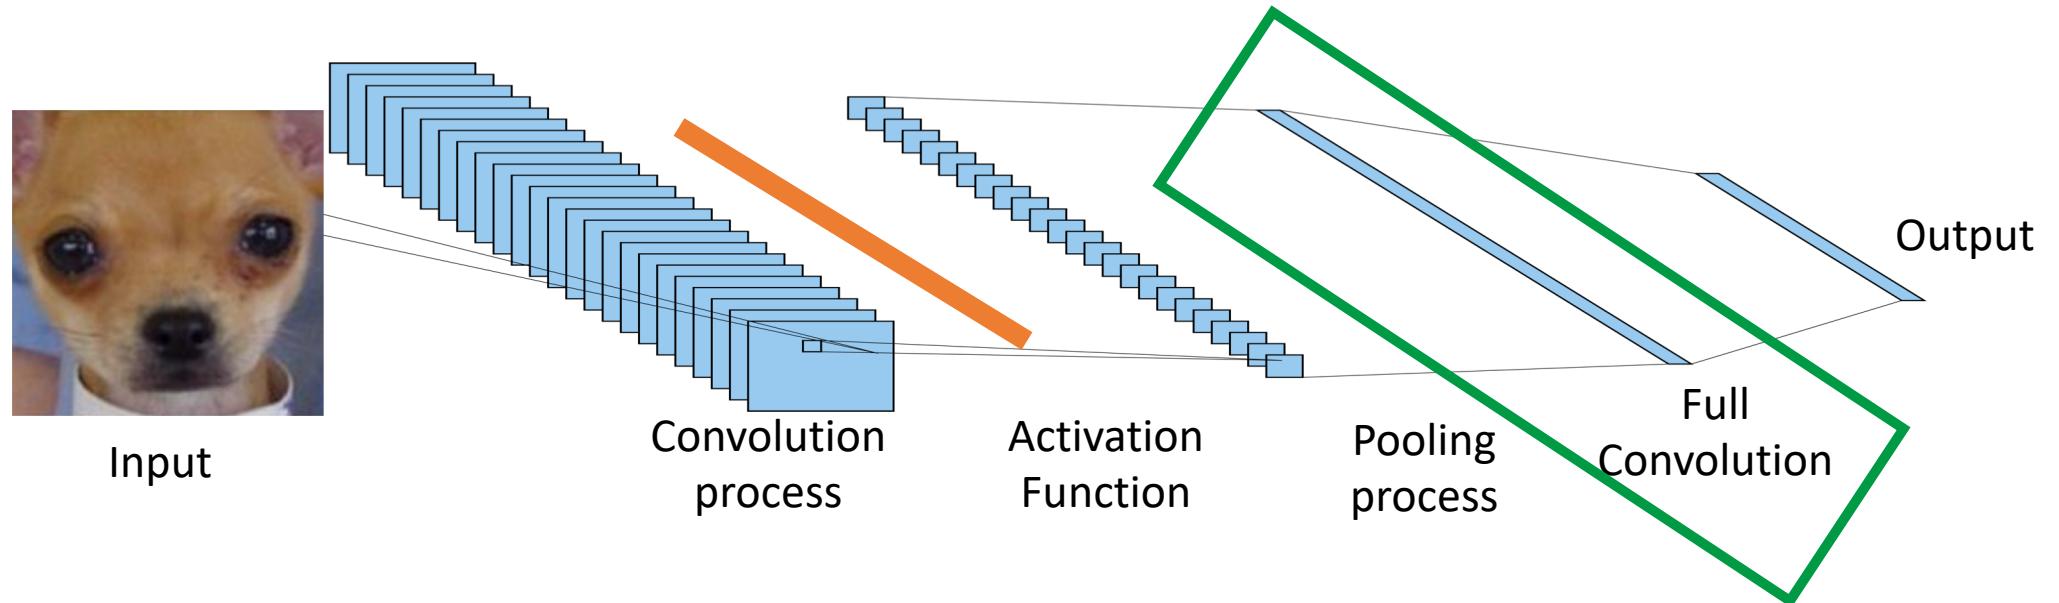


Pooling process

- Reduces the dimensionality of feature maps but retains important information.
 - Focuses on the main features
 - Reduce computational costs - gets rid of extra data
 - Reduce overfitting - combines neighboring values into a single one
- Types of pooling
 - **Average Pooling:** average value for each patch on the feature map.
 - **Maximum Pooling:** maximum value for each patch of the feature map.

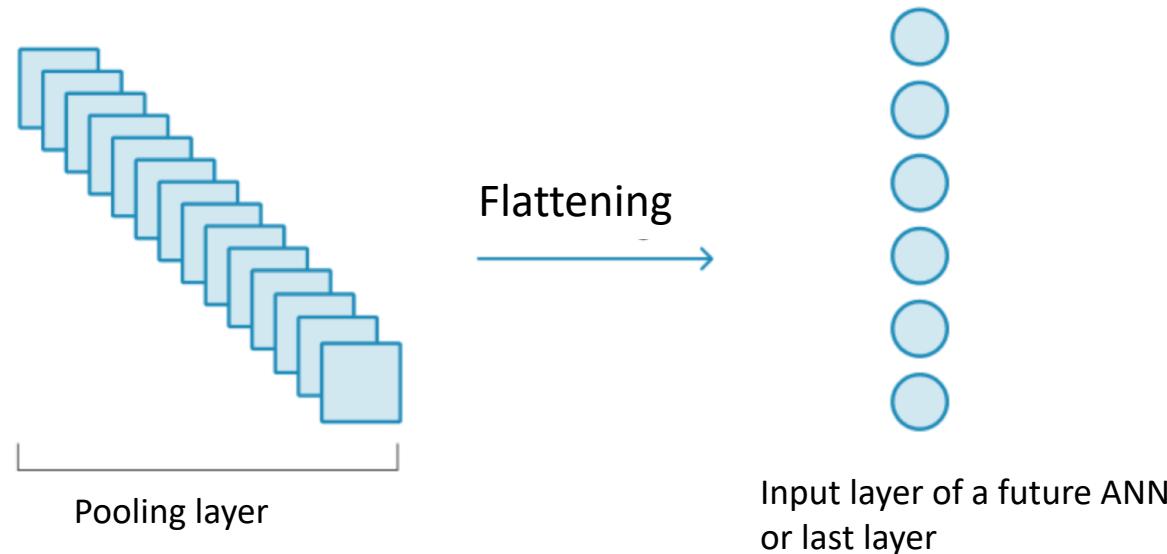


CNN Steps

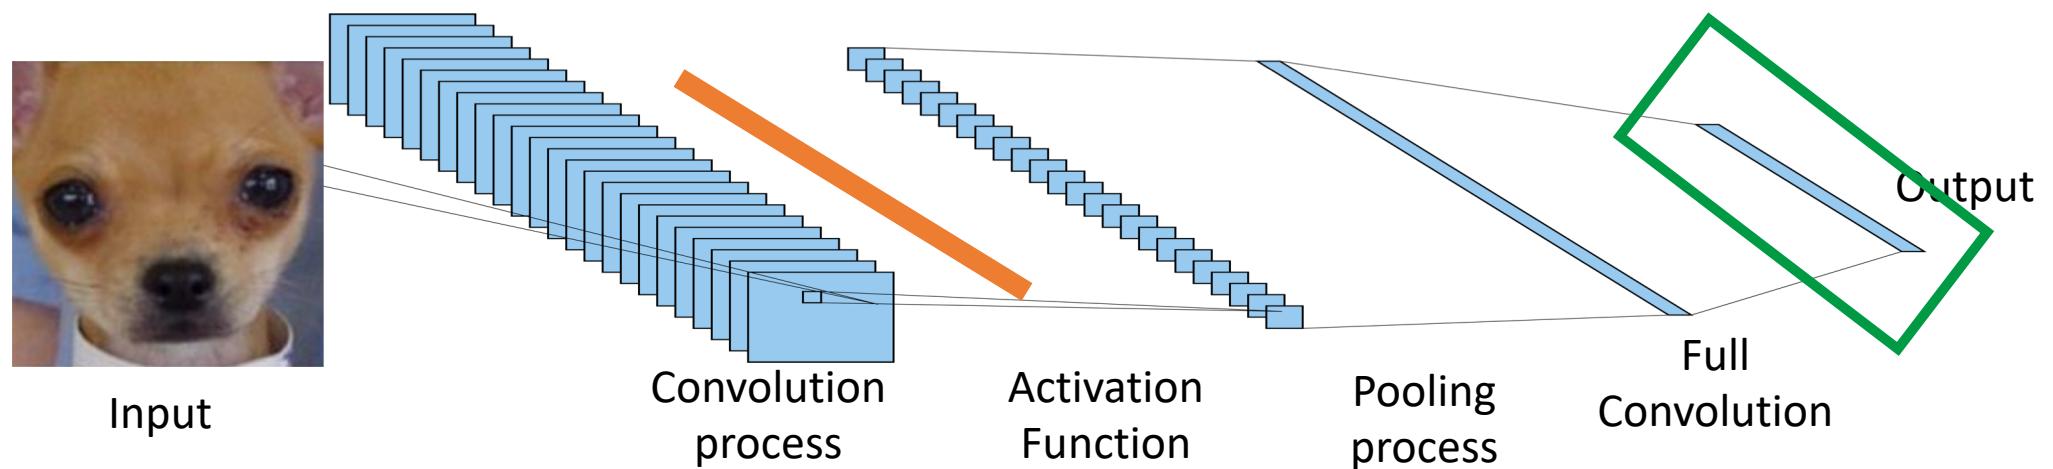


Full convolution

- Also known as Dense Layer
- Facilitated by a process called Flattening.
- Flattening transforms the feature maps into **one** vector
- The vector fed into a neural network of n layers.
- Last layer has c neurons, representing the **number of classes** to classify.

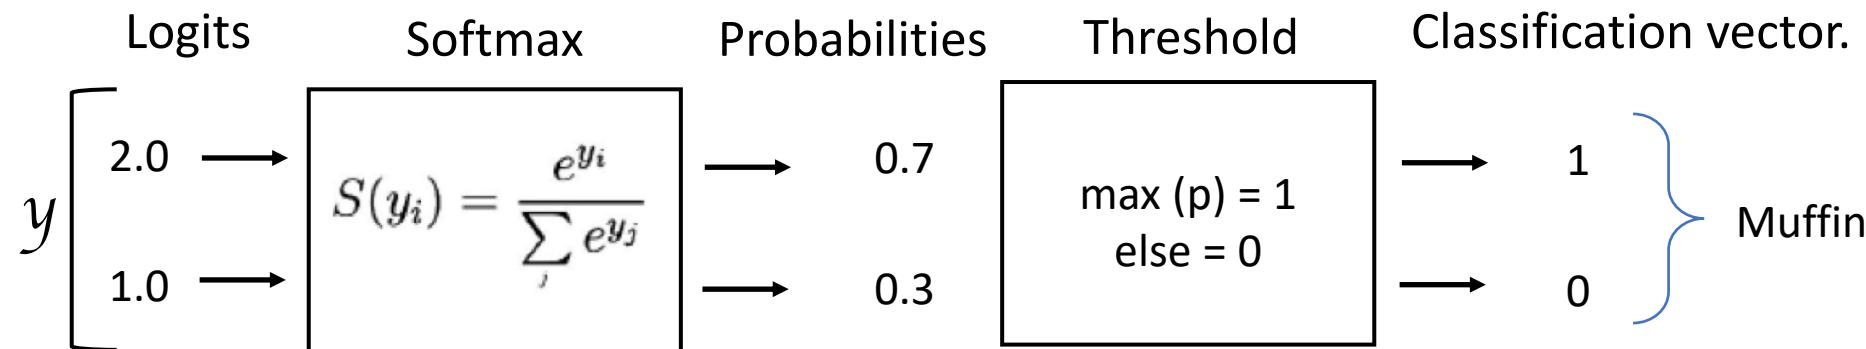


CNN Steps

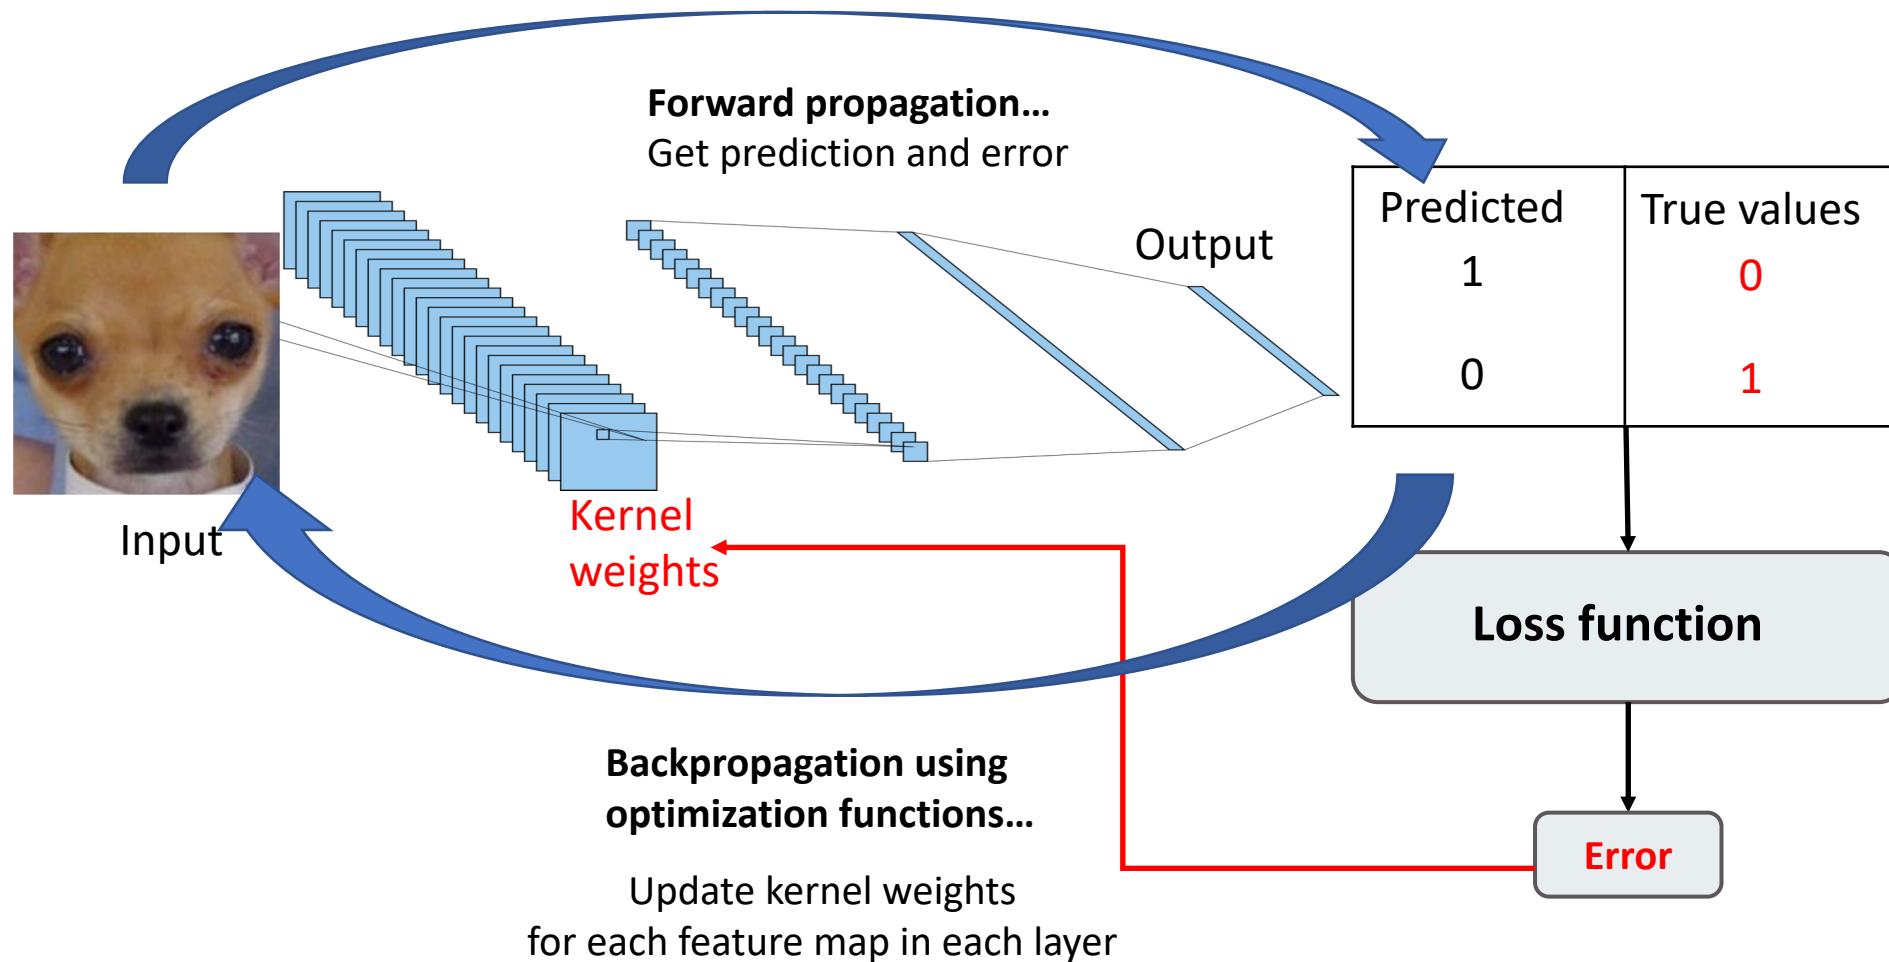


Output

- **Problem:** Last layer outputs numbers (a.k.a **logits**) are hard to interpret.
- **Solution:** normalize these numbers with an activation function
 - e.g. Softmax function results in probabilities with sum equal to 1
 - Threshold probabilities to generate a classification vector.



Training - updating the kernel weights



Overfitting

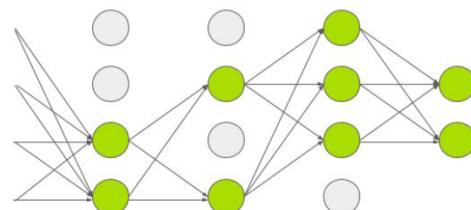
Algorithm performs well on the training set but performs poorly on the testing set => the algorithm is overfitted on the Training data.

Remedy: apply Regularization

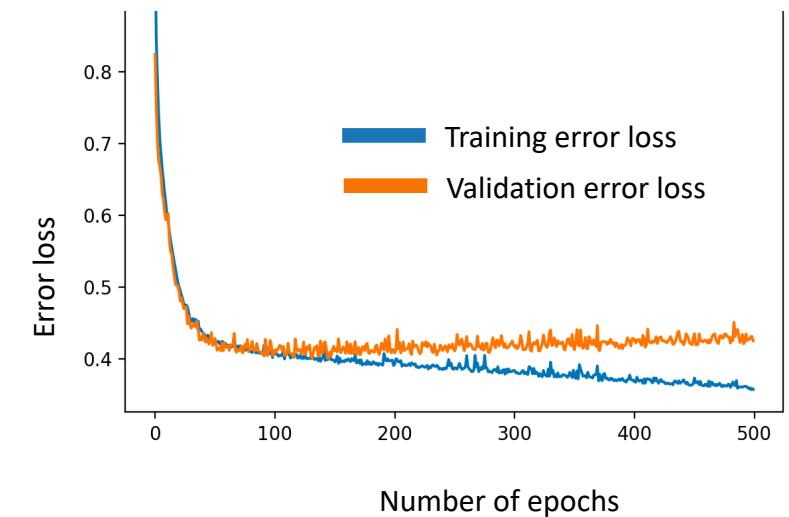
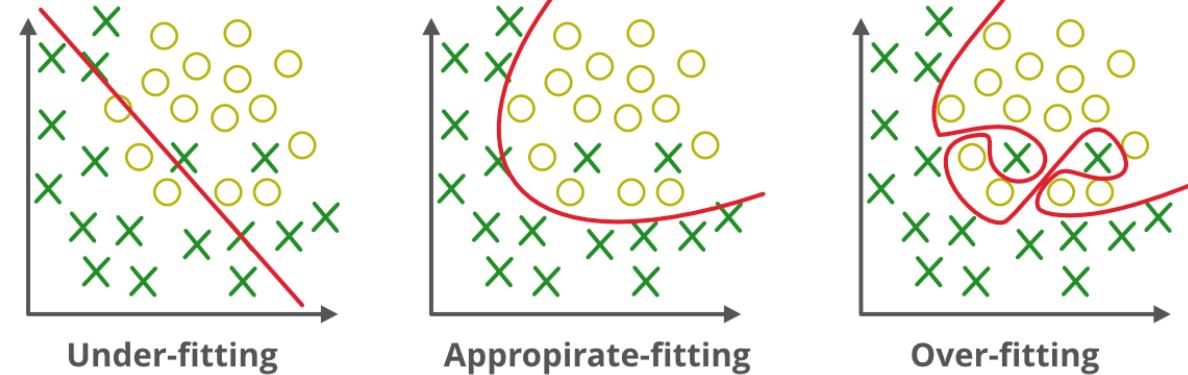
1. Adds penalty to weights and biases

- **L1** : the penalty is proportional to the absolute value of the weight coefficients.
- **L2** : the penalty is proportional to the square of the value of the weight coefficients

2. Drop out: randomly turns off a subset of neurons in a layer



3. Regularization rate: controls the amount of regularization applied

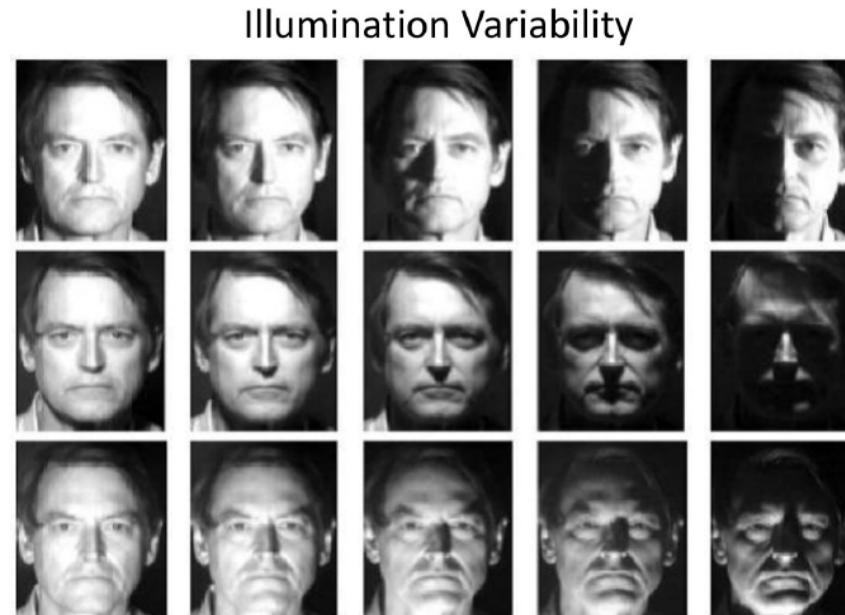
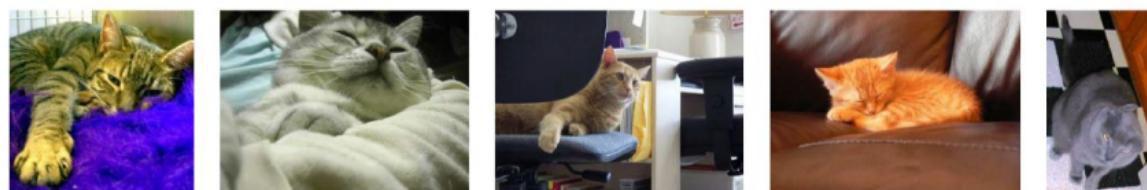


Computer Vision is Hard

Intra-Class Variability



Pose Variability

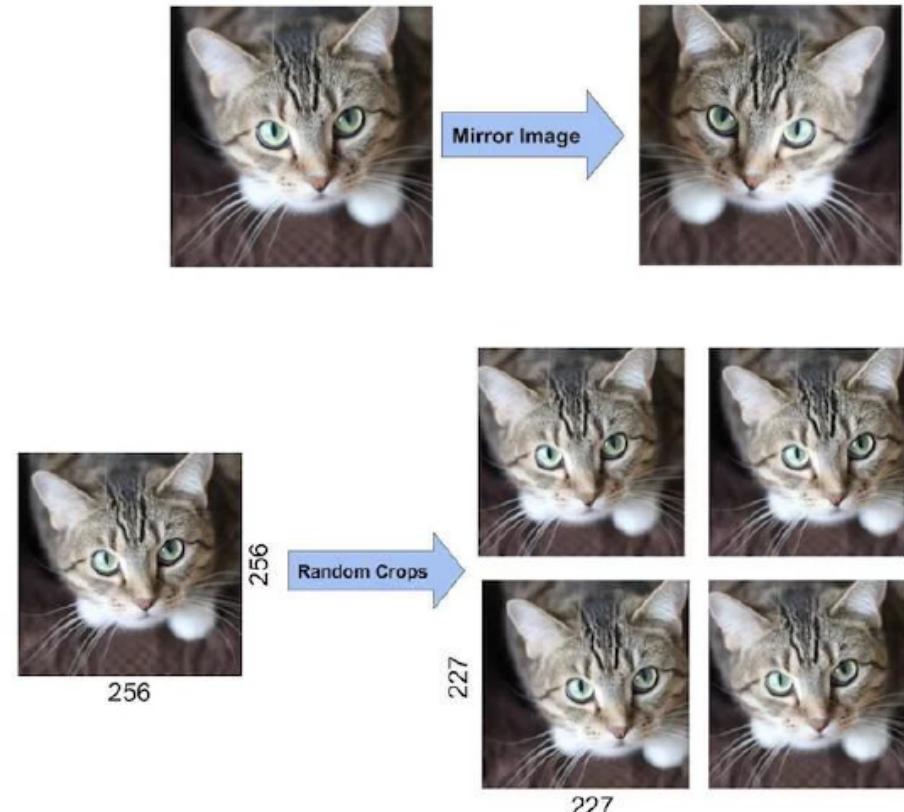


Occlusion



Solution (I): Data augmentation

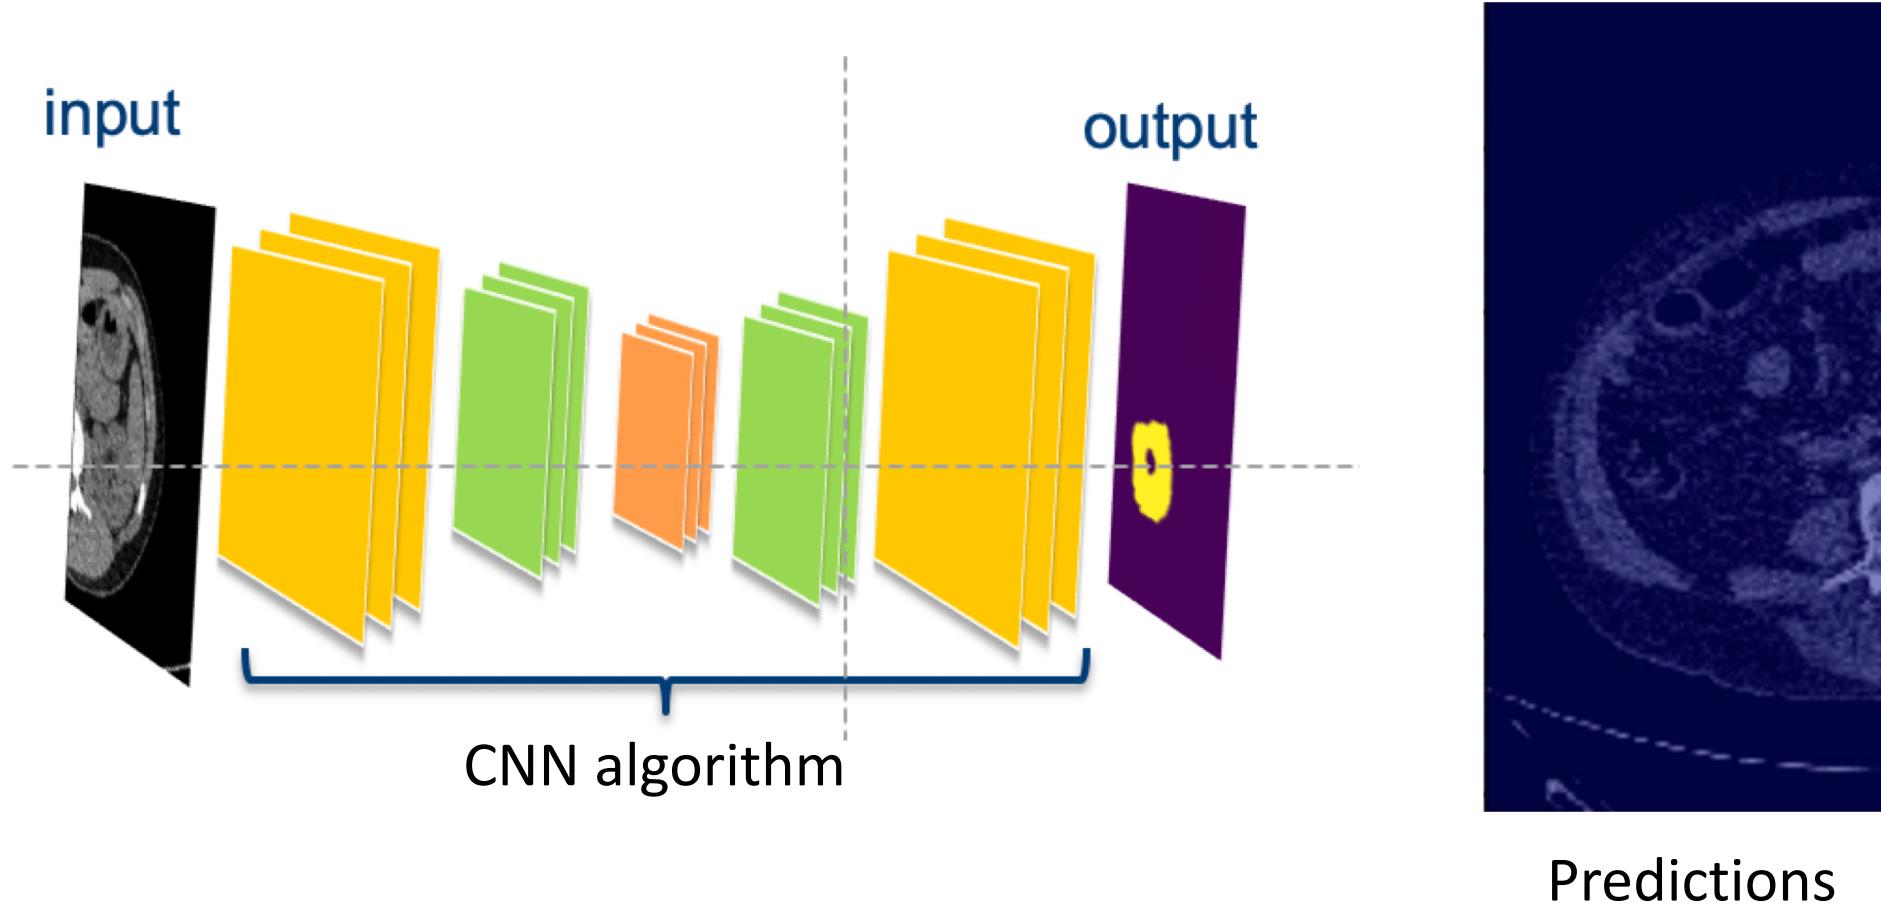
- Data augmentation is used as a way to make models more robust.
- It reduces overfitting by letting models learn harder and more general data representations.
- Most popular:
 - Noise addition
 - Mirroring
 - Rotating
 - Resizing
 - Cropping
 - Occlusion
- Newer methods:
 - Synthetic data



Solution (II) - Transfer Learning & Fine-Tuning

- Transfer learning is a machine learning method where a model developed for a task is reused as the starting point for a model on a second task.
- Start with a model that is pre-trained on a large amount of data for a task.
- Fine-tune some parts of the model to converge faster on a similar/relevant task where data isn't as abundant.
- Most popular is fine-tuning the feature extractor and adding a new classifier on top.

Application example – kidney segmentation



Deep Learning News and new use cases

- OpenAI: <https://openai.com/>
- Deepmind: <https://deepmind.com/blog>
- Nvidia: <https://blogs.nvidia.com/blog/category/deep-learning>
- Towardsdatascience: <https://towardsdatascience.com>

Emerging Tech & Research

Reddit News:

- <https://www.reddit.com/r/deeplearning/>
- <https://www.reddit.com/r/MachineLearning/>

Research:

- <https://distill.pub>
- <http://www.arxiv-sanity.com/>
- <https://paperswithcode.com/>

Very good Tutorials

Computer Vision: <https://www.pyimagesearch.com/>

Mixed: <https://machinelearningmastery.com/blog/>

Pretrained Tensorflow models

Tensorflow.keras.applications:

- <http://keras.io/applications/>
- https://www.tensorflow.org/versions/r2.0/api_docs/python/tf/keras/applications

Tensorflow Hub:

- <https://tfhub.dev/>
- <https://www.tensorflow.org/hub>

Tensorflow Model Collection:

- <https://github.com/tensorflow/models>

There are a lot more repositories

Hands on exercise

1. Open the following link

<https://drive.google.com/file/d/1TH-gRPnWKFsTi32uTh0JjeOp5iOohXfZ/view?usp=sharing>

2. Click “open with Google Colaboratory”, highlighted in red in the image below



3. Go to “File” at the left of the top menu, and click “Save a copy in Drive”
4. Click the icon at the top left menu
5. Open your copied file “Copy of CNN-intro.ipynb”